

## *Day 2:*

- GPU Structure *and* the Problem of Branch Divergence
- Implementation the reduction algorithm on GPU

Dr. Tuan-Tu Tran  
trant@uni-mainz.de




JOHANNES GUTENBERG  
UNIVERSITÄT MAINZ

# Part 1: GPU Structure *and* the Problem of Branch Divergence

# Why multi/many-core?

## Computing Power:

- Number of transistors on a chip: has doubled every two years, since 1965
- Frequency : has not increased anymore, since the beginning of 2000s
- Performance increased by **integrating multiple “cores” on a chip.**



(image source: [3])

# Multicore and Manycore

- Meaning: multi- ~ many-
- Actually:
  - Multicore CPU: 2 cores, 4 cores, 8 cores
  - Manycore GPU: 100 to 1000 cores

?

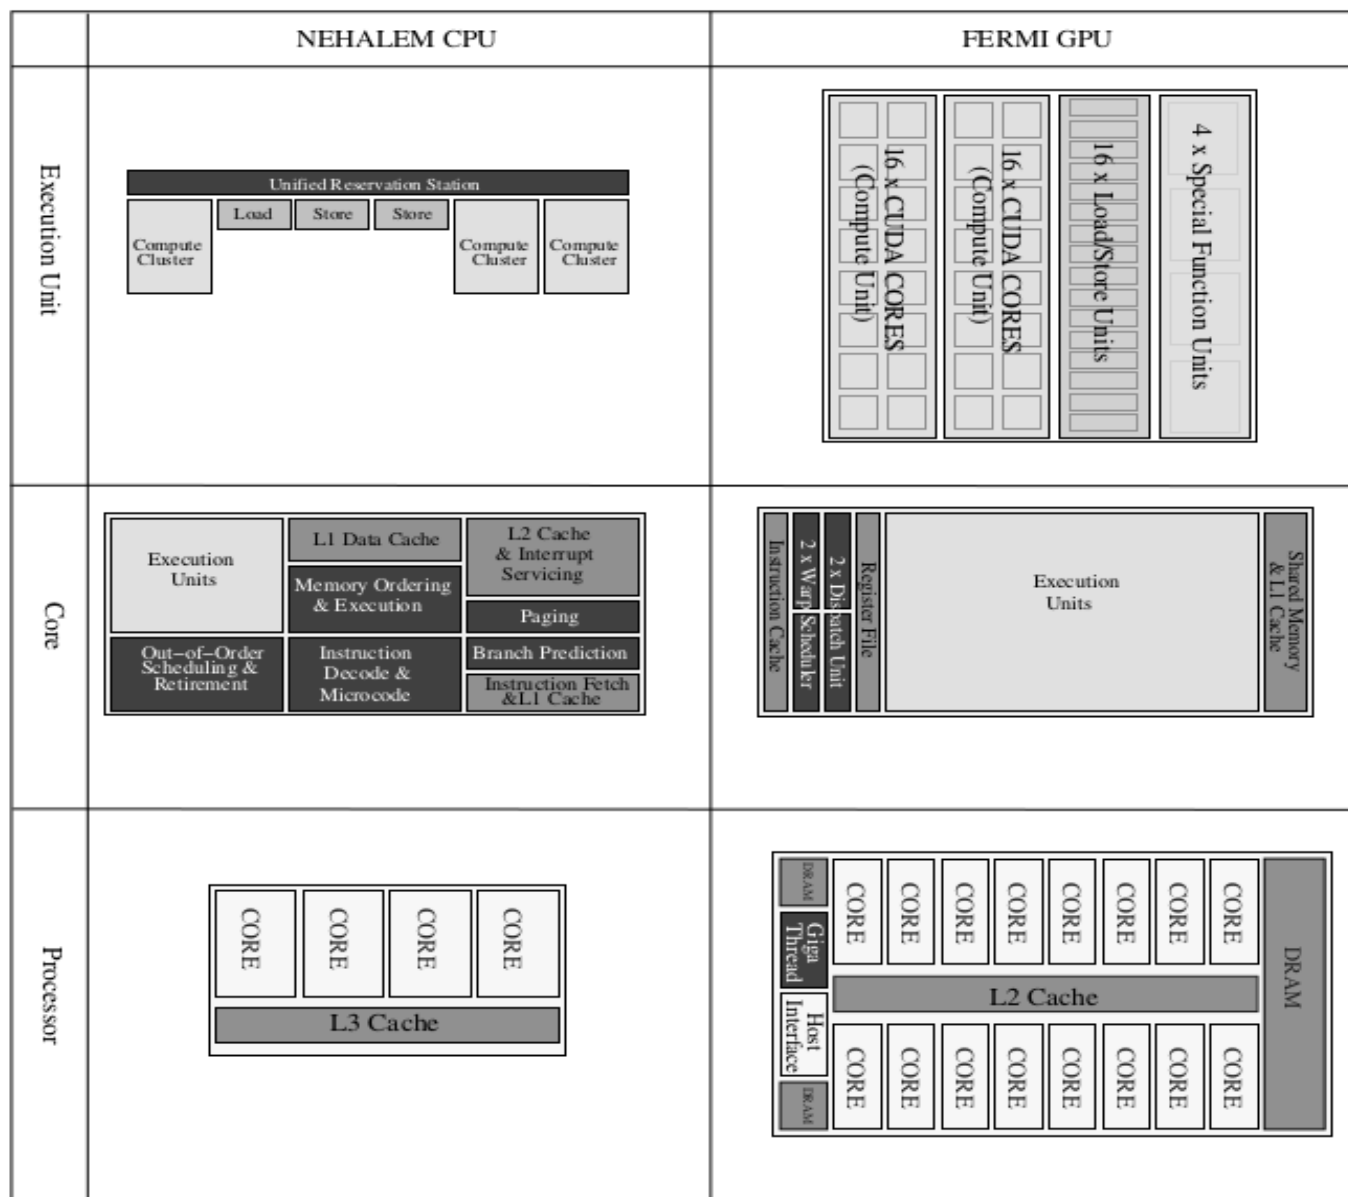
# Multicore and Manycore

- Meaning: multi- ~ many-
- Actually:
  - Multicore CPU: 2 cores, 4 cores, 8 cores
  - Manycore GPU: 100 to 1000 cores

The GPU cores has


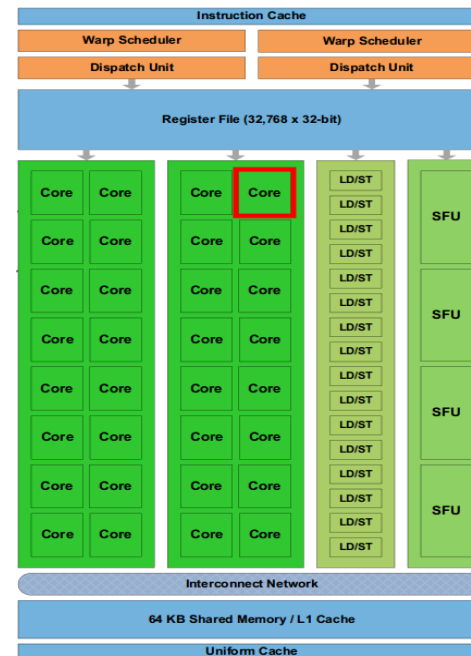

- *lower frequency*
- *simpler structure*

than the CPU cores



(Architecture comparison between Intel NEHALEM CPU and NVIDIA Fermi GPU at three levels: processor, core and execution unit, based on [5] and [6])

# Example: Inside a Fermi GPU

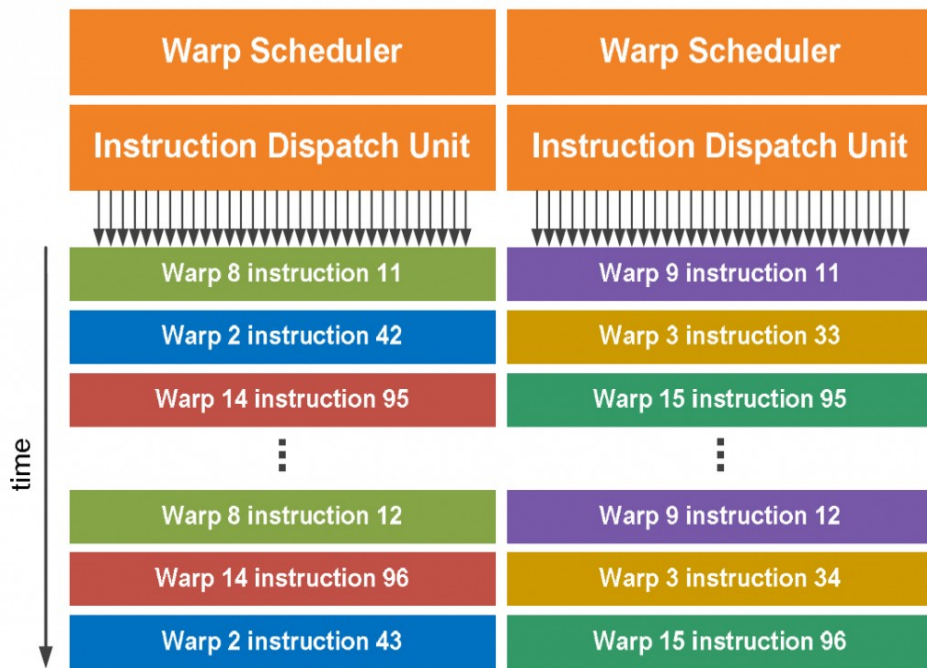


(images source: [6])

- 1 Fermi GPU contains 16 Streaming Multiprocessors (SM)
- 1 SM contains 32 CUDA cores

# CUDA Warp

32 consecutive threads = 1 warp



(image source: [6])

- Program ~ Sequence of Instructions  
(In this practical, a kernel can be considered as a program)
  - SPMD: Single Program Multiple Data
  - SIMD: Single Instruction Multiple Data
- Between the threads in the same warp: SIMD
- Between the threads in different warp: SPMD



# Helloworld



JOHANNES GUTENBERG  
UNIVERSITÄT MAINZ

**./helloworld 1 128**

**blocksPerGrid = 1, threadsPerBlock = 128**

Hello, I am thread 64, of block 0  
Hello, I am thread 65, of block 0  
Hello, I am thread 66, of block 0  
Hello, I am thread 67, of block 0  
Hello, I am thread 68, of block 0  
Hello, I am thread 69, of block 0  
Hello, I am thread 70, of block 0  
Hello, I am thread 71, of block 0  
Hello, I am thread 72, of block 0  
Hello, I am thread 73, of block 0  
Hello, I am thread 74, of block 0  
Hello, I am thread 75, of block 0  
Hello, I am thread 76, of block 0  
Hello, I am thread 77, of block 0  
Hello, I am thread 78, of block 0  
Hello, I am thread 79, of block 0  
Hello, I am thread 80, of block 0  
Hello, I am thread 81, of block 0  
Hello, I am thread 82, of block 0  
Hello, I am thread 83, of block 0  
Hello, I am thread 84, of block 0  
Hello, I am thread 85, of block 0  
Hello, I am thread 86, of block 0  
Hello, I am thread 87, of block 0  
Hello, I am thread 88, of block 0  
Hello, I am thread 89, of block 0  
Hello, I am thread 90, of block 0  
Hello, I am thread 91, of block 0  
Hello, I am thread 92, of block 0  
Hello, I am thread 93, of block 0  
Hello, I am thread 94, of block 0  
Hello, I am thread 95, of block 0

warp 2

Hello, I am thread 96, of block 0  
Hello, I am thread 97, of block 0  
Hello, I am thread 98, of block 0  
Hello, I am thread 99, of block 0  
Hello, I am thread 100, of block 0  
Hello, I am thread 101, of block 0  
Hello, I am thread 102, of block 0  
Hello, I am thread 103, of block 0  
Hello, I am thread 104, of block 0  
Hello, I am thread 105, of block 0  
Hello, I am thread 106, of block 0  
Hello, I am thread 107, of block 0  
Hello, I am thread 108, of block 0  
Hello, I am thread 109, of block 0  
Hello, I am thread 110, of block 0  
Hello, I am thread 111, of block 0  
Hello, I am thread 112, of block 0  
Hello, I am thread 113, of block 0  
Hello, I am thread 114, of block 0  
Hello, I am thread 115, of block 0  
Hello, I am thread 116, of block 0  
Hello, I am thread 117, of block 0  
Hello, I am thread 118, of block 0  
Hello, I am thread 119, of block 0  
Hello, I am thread 120, of block 0  
Hello, I am thread 121, of block 0  
Hello, I am thread 122, of block 0  
Hello, I am thread 123, of block 0  
Hello, I am thread 124, of block 0  
Hello, I am thread 125, of block 0  
Hello, I am thread 126, of block 0  
Hello, I am thread 127, of block 0

warp 3

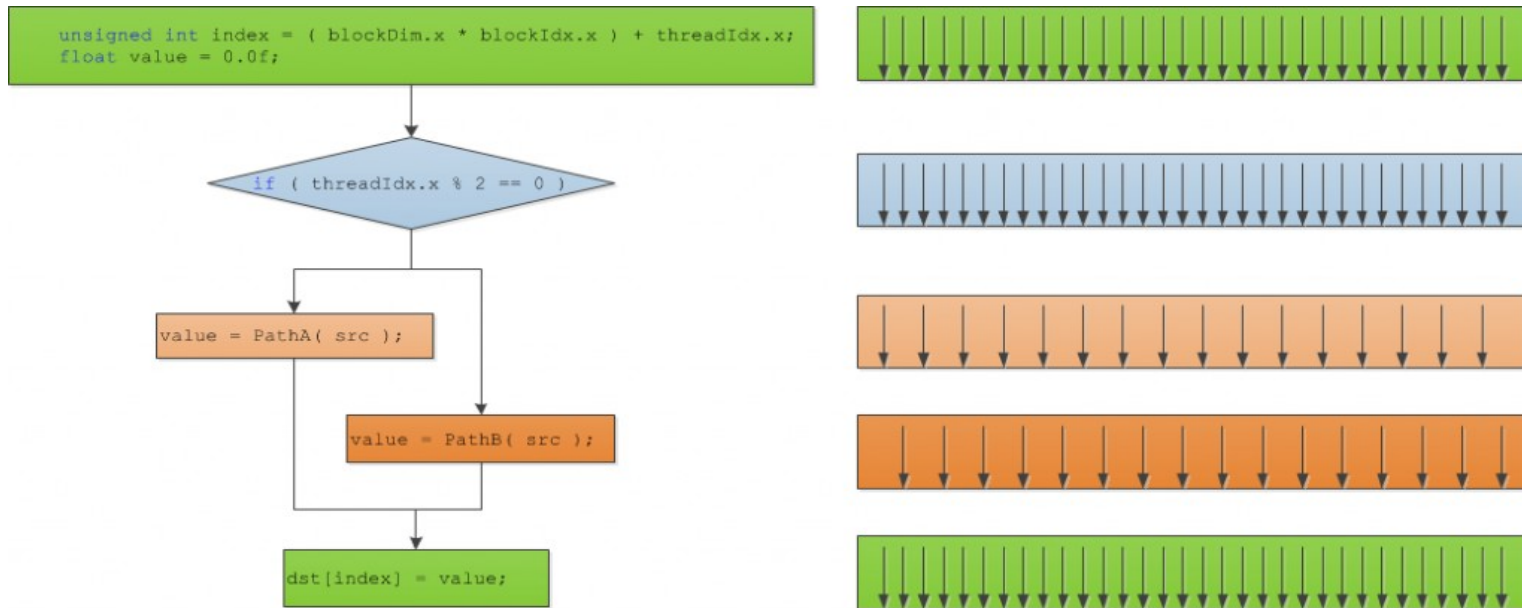
Hello, I am thread 0, of block 0  
Hello, I am thread 1, of block 0  
Hello, I am thread 2, of block 0  
Hello, I am thread 3, of block 0  
Hello, I am thread 4, of block 0  
Hello, I am thread 5, of block 0  
Hello, I am thread 6, of block 0  
Hello, I am thread 7, of block 0  
Hello, I am thread 8, of block 0  
Hello, I am thread 9, of block 0  
Hello, I am thread 10, of block 0  
Hello, I am thread 11, of block 0  
Hello, I am thread 12, of block 0  
Hello, I am thread 13, of block 0  
Hello, I am thread 14, of block 0  
Hello, I am thread 15, of block 0  
Hello, I am thread 16, of block 0  
Hello, I am thread 17, of block 0  
Hello, I am thread 18, of block 0  
Hello, I am thread 19, of block 0  
Hello, I am thread 20, of block 0  
Hello, I am thread 21, of block 0  
Hello, I am thread 22, of block 0  
Hello, I am thread 23, of block 0  
Hello, I am thread 24, of block 0  
Hello, I am thread 25, of block 0  
Hello, I am thread 26, of block 0  
Hello, I am thread 27, of block 0  
Hello, I am thread 28, of block 0  
Hello, I am thread 29, of block 0  
Hello, I am thread 30, of block 0  
Hello, I am thread 31, of block 0

warp 0

Hello, I am thread 32, of block 0  
Hello, I am thread 33, of block 0  
Hello, I am thread 34, of block 0  
Hello, I am thread 35, of block 0  
Hello, I am thread 36, of block 0  
Hello, I am thread 37, of block 0  
Hello, I am thread 38, of block 0  
Hello, I am thread 39, of block 0  
Hello, I am thread 40, of block 0  
Hello, I am thread 41, of block 0  
Hello, I am thread 42, of block 0  
Hello, I am thread 43, of block 0  
Hello, I am thread 44, of block 0  
Hello, I am thread 45, of block 0  
Hello, I am thread 46, of block 0  
Hello, I am thread 47, of block 0  
Hello, I am thread 48, of block 0  
Hello, I am thread 49, of block 0  
Hello, I am thread 50, of block 0  
Hello, I am thread 51, of block 0  
Hello, I am thread 52, of block 0  
Hello, I am thread 53, of block 0  
Hello, I am thread 54, of block 0  
Hello, I am thread 55, of block 0  
Hello, I am thread 56, of block 0  
Hello, I am thread 57, of block 0  
Hello, I am thread 58, of block 0  
Hello, I am thread 59, of block 0  
Hello, I am thread 60, of block 0  
Hello, I am thread 61, of block 0  
Hello, I am thread 62, of block 0  
Hello, I am thread 63, of block 0

warp 1

# Problem of Branch Divergence



(image source: [4])

- Increase the running time of the whole program
- Waste computing resources

## *Part 2:*

# Implementation the reduction algorithm on GPU


# Reduction



JOHANNES GUTENBERG  
UNIVERSITÄT MAINZ

- Given an array of value  $\Rightarrow$  reduce them to a single value.
  - Sum reduction  $\Rightarrow$  sum of all values in the array
  - Max reduction  $\Rightarrow$  maximal value of the array
  - etc
- In this practical, we work with the Sum Reduction

# Parallel reduction algorithm



*(image source: [1])*

# Implementation Manual

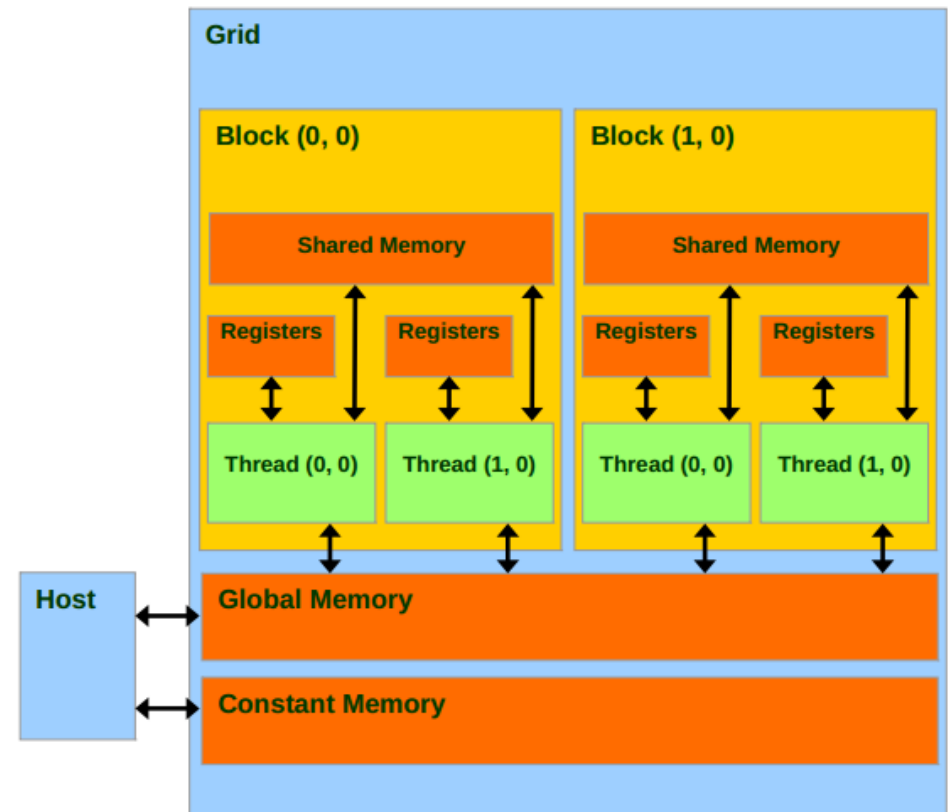
- Based on the given skeleton: `cudaReduction.cu` (or from the source code of the `squareArray` application in Day 1)
- Using the ideas and the pseudo-codes in:  
[1] Mark Harris, Optimizing Parallel Reduction in CUDA,  
[http://developer.download.nvidia.com/compute/cuda/1.1-Beta/x86\\_website/projects/reduction/doc/reduction.pdf](http://developer.download.nvidia.com/compute/cuda/1.1-Beta/x86_website/projects/reduction/doc/reduction.pdf)
- **Basic implementation:** the tree-based approach with 1 thread block, with small data (up to 1024 elements), by using 2 algorithms:
  - Interleaved addressing with divergent branching [1, page 7]
  - Interleaved addressing without divergent branching [1, page 11]
- **Advanced implementation:** based on the basic implementation for small data, write the full CUDA application to process the large data of any size, by using recursive kernel invocation [1, page 7]. We suppose that the size of the input array is the power of 2, and up to 32M (33554432 elements).
- **Extend implementation:** the tree-based approach with 1 thread block, with small data (up to 1024 elements), by using the other ideas and optimizations in [1],

# A brief introduction to shared memory



JOHANNES GUTENBERG  
UNIVERSITÄT MAINZ

- Much faster than global memory
- Small (upto 48 KB)
- Shared among the thread in the same block
- *(more details will be discussed in Day 3)*



(image credit: NVIDIA)

# Using the shared memory inside the CUDA code

The shared memory can be declared and allocated **inside** or **outside** the kernel:

Inside the kernel, the size must be fixed

```
#define LEN 8
__global__ kernell(...)
{
    __shared__ int  arr1[LEN];
    __shared__ int  arr2[16];
    ...
}
```

Outside the kernel, the size can be set in the host code, by using the third configuration argument when invoking the kernel:

```
#define LEN 8
__global__ kernell(...)
{
    extern __shared__ int  arr[];
    ...
}
//
int main(int argc, char** argv)
{
    ...
    n = 24;
    sharedSize = n * sizeof(int);
    kernell<<<nBlock,nThread,sharedSize>>>(...)
    ...
}
```



# Declare multiple `__shared__` memory segments



JOHANNES GUTENBERG  
UNIVERSITÄT MAINZ

```
__global__ kernel1(..., len1, len2)
{
    extern __shared__ int arr[];
    int* arr1 = arr;
    int* arr2 = arr1 + len1;
    ...
}

//
int main(int argc, char** argv)
{
    ...
    len1 = 8;
    len2 = 16;
    n = len1 + len2;
    sharedSize = n * sizeof(int);
    kernel1<<<nBlock, nThread, sharedSize>>>(..., len1, len2)
    ...
}
```

# Thread Synchronization

- Using the command: `__syncthreads()`
- Example (from [1]):

## Reduction #1: Interleaved Addressing

```
__global__ void reduce0(int *g_idata, int *g_odata) {
    extern __shared__ int sdata[];

    // each thread loads one element from global to shared mem
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
    sdata[tid] = g_idata[i];
    __syncthreads();

    // do reduction in shared mem
    for(unsigned int s=1; s < blockDim.x; s *= 2) {
        if (tid % (2*s) == 0) {
            sdata[tid] += sdata[tid + s];
        }
        __syncthreads();
    }

    // write result for this block to global mem
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

# Reference

[1] Mark Harris, Optimizing Parallel Reduction in CUDA,

[http://developer.download.nvidia.com/compute/cuda/1.1-Beta/x86\\_website/projects/reduction/doc/reduction.pdf](http://developer.download.nvidia.com/compute/cuda/1.1-Beta/x86_website/projects/reduction/doc/reduction.pdf)

[2] Hendrik Lensch and Robert Strzodka, Massively Parallel Computing with Cuda, Control Flow

<http://www.mpi-inf.mpg.de/~strzodka/lectures/ParCo08/slides/Par03-ControlFlow.pdf>

[3] Shalf, J., Asanovic, K., Patterson, D., Keutzer, K., Mattson, T., and Yelick, K. (2009). The Manycore Revolution: Will HPC Lead or Follow? SciDAC Review, 14.

[4] Jeremiah van Oosten, Optimizing CUDA Application, <http://3dgep.com/?p=2081>

[5] Semin, A. (2009). Inside Intel Nehalem Microarchitecture. In NOTUR 2009, The 8<sup>th</sup> Annual Meeting on High Performance Computing and Infrastructure in Norway.

[6] NVIDIA Corp. NVIDIA's Next Generation CUDA Compute Architecture: Fermi. white paper, version 1.1

And some images from the CUDA technical documents of NVIDIA



JOHANNES GUTENBERG  
UNIVERSITÄT MAINZ

# Question ?