

Day3

- Discussion on the parallel reduction algorithm
- Implementation the parallel matrix-matrix multiplication algorithm on GPU

Dr. Tuan-Tu Tran
trant@uni-mainz.de

Part 1

Discussion on the parallel reduction algorithm

[1] Mark Harris, Optimizing Parallel Reduction in CUDA,
http://developer.download.nvidia.com/compute/cuda/1.1-Beta/x86_website/projects/reduction/doc/reduction.pdf

Ver-1: the branch divergence problem

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    if (tid % (2*s) == 0) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

`N = 64, threadsPerBlock = 64 (blockDim.x = 64)`

“active” threads in each loop:

`S = 1 => s*2 = 2:`
threads: 0, 2, 4, ... 30, 32, ... 62

`S = 2 => s*2 = 4:`
Threads: 0, 4, 8 ... 28, 32, ... 60

...

`S = 16 => s*2 = 32`
Threads: 0, 32

`S = 32 => s*2 = 64`
Threads: 0

There is always the branch divergence problem

Ver-2: reduce the branch divergence problem



```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    int index = 2 * s * tid;  
  
    if (index < blockDim.x) {  
        sdata[index] += sdata[index + s];  
    }  
    __syncthreads();  
}
```

No branch divergence

The branch divergence problem in the last 5 steps (Loop2 to Loop6)

$N = 64$, threadsPerBlock = 64 (blockDim.x = 64)

“active” threads in each loop:

Loop1: $s = 1 \Rightarrow s*2 = 2$:
tid = 0: \Rightarrow sData[0], sData[1]
tid = 1: \Rightarrow sData[2], sData[3]
...
tid = 16: \Rightarrow sData[32], sData[33]
...
tid = 31: \Rightarrow sData[62], sData[63]
Loop2: $s = 2 \Rightarrow s*2 = 4$:
tid = 0: \Rightarrow sData[0], sData[2]
tid = 1: \Rightarrow sData[4], sData[6]
...
tid = 16: \Rightarrow index = 64 \Rightarrow not active
...
Loop5: $S = 16 \Rightarrow s*2 = 32$
tid = 0 \Rightarrow sData[0], sData[16]
tid = 1 \Rightarrow sData[32], sData[48]
tid = 2 \Rightarrow not active
Loop6: $S = 32 \Rightarrow s*2 = 64$
tid = 0 \Rightarrow sData[0], sData[32]
tid = 1 \Rightarrow not active

Ver-2: The bank conflict problem



```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    int index = 2 * s * tid;  
  
    if (index < blockDim.x) {  
        sdata[index] += sdata[index + s];  
    }  
    __syncthreads();  
}
```

Loop1: Bank conflict between:
tid0 and tid16
tid1 and tid17
...

Loop2: Bank conflict between:
tid0 and tid8
tid1 and tid9
...

$N = 64$, threadsPerBlock = 64 (blockDim.x = 64)

“active” threads in each loop:

Loop1: $s = 1 \Rightarrow s*2 = 2$:

tid = 0: \Rightarrow sData[0], sData[1]

tid = 1: \Rightarrow sData[2], sData[3]

...

tid = 16: \Rightarrow sData[32], sData[33]

...

tid = 31: \Rightarrow sData[62], sData[63]

Loop2: $s = 2 \Rightarrow s*2 = 4$:

tid = 0: \Rightarrow sData[0], sData[2]

tid = 1: \Rightarrow sData[4], sData[6]

tid = 8: \Rightarrow sData[32], sData[34]

...

tid = 16: \Rightarrow index = 64 \Rightarrow not active

...

Loop5: $S = 16 \Rightarrow s*2 = 32$

tid = 0 \Rightarrow sData[0], sData[16]

tid = 1 \Rightarrow sData[32], sData[48]

tid = 2 \Rightarrow not active

Loop6: $S = 32 \Rightarrow s*2 = 64$

tid = 0 \Rightarrow sData[0], sData[32]

tid = 1 \Rightarrow not active

Ver-3: Free of bank conflict problem



JOHANNES GUTENBERG
UNIVERSITÄT MAINZ

```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {  
    if (tid < s) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

$N = 64$, threadsPerBlock = 64 (blockDim.x = 64)

“active” threads in each loop:

Loop1: $s = 32$:

tid = 0: \Rightarrow sData[0], sData[32]

tid = 1: \Rightarrow sData[1], sData[33]

...

tid = 16: \Rightarrow sData[16], sData[48]

...

tid = 31: \Rightarrow sData[31], sData[63]

Loop2: $s = 16$:

tid = 0: \Rightarrow sData[0], sData[16]

tid = 1: \Rightarrow sData[1], sData[17]

...

tid = 15: \Rightarrow sData[15], sData[31]

tid = 16: \Rightarrow not active

...

Loop5: $S = 2$:

tid = 0 \Rightarrow sData[0], sData[2]

tid = 1 \Rightarrow sData[1], sData[3]

tid = 2 \Rightarrow not active

Loop6: $S = 1 \Rightarrow s*2 = 64$

tid = 0 \Rightarrow sData[0], sData[1]

tid = 1 \Rightarrow not active

Ver-1: the branch divergence problem

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    if (tid % (2*s) == 0) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

$N = 64$, threadsPerBlock = 64 (blockDim.x = 64)

“active” threads in each loop:

$S = 1 \Rightarrow s*2 = 2$:

threads: 0, 2, 4, ... 30, 32, ... 62

$S = 2 \Rightarrow s*2 = 4$:

Threads: 0, 4, 8 ... 28, 32, ... 60

...

$S = 16 \Rightarrow s*2 = 32$

Threads: 0, 32

$S = 32 \Rightarrow s*2 = 64$

Threads: 0

Free of the bank conflict problem

Full application

- Solution: `cudaReduction.cu`

(thanks to Tassilo Kugelstadt and Denise Scherzinger)

- Discussion:
 - Padding to the nearest power of 2
 - Recursive loop in the host
 - The output array in the GPU

Part 2

Implementation the parallel matrix-matrix multiplication algorithm on GPU

Parallel matrix-matrix multiplication algorithm on GPU

$$C_{mn} = A_{mp} \times B_{pn}$$

- Using 2-Dimension Block and Grid
- In this practical we use **square matrices of integer, where $m = n = p = N$** . You can use the same skeleton of the reduction application:
 - Assume that $A = B$ (read data only 1 time)
 - Read N^2 integer, instead of N
- N is the power of 2, and up to $2^{12} = 4K$ (4096).
- **Basic Implementations** (either Case 1 or Case 2)
 - 1 block, small data ($N = 8$)
 - Directly access to the global memory
 - Using shared memory
- **Advanced Implementation** (either Case 1 or Case 2)
 - Multiple block, large data
 - Using shared memory and apply optimizations presented in [1], from page 27 to 68

[2]. **Hendrik Lensch and Robert Strzodka**, Massively Parallel Computing With CUDA: Memory Access, <http://www.mpi-inf.mpg.de/~strzodka/lectures/ParCo08/slides/Par02-Memory.pdf>



JOHANNES GUTENBERG
UNIVERSITÄT MAINZ

Question ?